

MatSuite Writeup

George Saussy

October 18, 2016

In this project we attempt to build a GPU optimized implementation of Krylov wave propagation for quantum states. In the process we, implement a small suite of functions designed to handle matrices. This system will be used to train an neural network that will speed up the calculation of the propagation of quantum states. The neural network can then be used to aid in the design of molecules which would be computationally infeasible otherwise.

1 Motivation

Physicist must frequently calculate the propagation of a quantum state. This calculation is often extremely computationally expensive, even for simple systems. This problem is only exasperated when many bodies are introduced or the Hamiltonian is not well behaved.

However, as physical methods play a greater role in chemistry, more efficient means of calculating wave propagation in many body systems must be explored. The aim of this research should be to bring down the cost of understanding the the behaviour of particles and molecules in generic settings, so that the methods found can be generally applied without needing to be customized for a given problem.

In the interest of generality the most generic form of a quantum system is given by the Schrodinger equation, $i\hbar d_t|\phi\rangle = H|\phi\rangle$, where $|\phi\rangle$ is the quantum state and H is the Hamiltonian under which the system is evolving. The generic solution to this equation is then

$$|\phi(t)\rangle = \exp(-iHt/\hbar)|\phi_0\rangle$$

However, this computation in practice is very expensive because of the exponential. In particular the generic definition of $\exp(A)$ for matrix A is given by

$$\exp(A) = \sum_{k=0}^{\infty} A^k/k!$$

which does not converge quickly for $\|A\| \gg 1$. (In fact, the error term given only a constant number of terms are used in the above equation will grow exponentially with $\|A\|$.) Bringing the computation time for matrix exponentiation down would allow theorists to consider systems with a greater set of pure quantum states, and potentially to numerically consider the behaviour of systems that were previously infeasible to treat.

In this project, I present the design and implementation of a suite of GPU optimized software to speed up these computations.

2 Mathematical background

The literature on efficient matrix exponentiation is fairly concise. There are fewer than a dozen papers that contribute to the field and only three or four that are core to the progress of the field. We choose to focus on two algorithms, the Pade approximation and the Krylov approximation.

The Pade approximation for matrix exponentiation is based on the Pade approximate, where a function is calculated by taking the ratio of two power series. We know the k -th polynomial coefficient in the Taylor expansion of $\exp(x)$ at zero is $k!^{-1}$. If we posit that

$$\exp(x) \approx \frac{P(x)}{Q(x)}$$

for

$$P(x) = \sum_{k=0}^p a_k x^k, Q(x) = \sum_{k=0}^q b_k x^k$$

then we can calculate the coefficients of the above approximation by equating terms in the approximation to terms in the Taylor expansion. For a given value of p and q the coefficients can be precomputed and stored in a lookup table for faster computation. In addition, algorithms exist to bound the error from this approximation, namely the one found in Ward 1977 ('Numerical Computation of the Matrix Exponential with Accuracy Estimate' SIAM J. NUMBER. ANAL. Vol 14, No. 4). This is essential to making sure the use of this approximation possible, because without an knowledge of the error the approximation is worthless.

The second approximation used is the Krylov approximation. The state-of-the-art implementation is [Expokit](#) by Roger Sidje. While this repository is stable, it is implemented in Fortran and Matlab. A C/SPIR-V implementation should outperform this suite. The Krylov approximation used in this project is an iterative method, which is optimized for sparse matrix exponentiation. In the generic case we use the Arnoldi iteration. (Later, Lanczos iteration should be implemented for the case the matrix can be guaranteed to be a valid Hamiltonian.)

For sparse matrix A and vector v , Arnoldi iteration calculates $\exp(A)v$ by finding an orthonormal basis for the set of vectors $\{v, Av, A^2v, \dots, A^m v\}$ for some typically small $m < \text{Dim}(v)$. This can be done efficiently by the Gram-Schmidt procedure. The algorithm then projects A onto this so-called Krylov basis. Where the eigenvalues can be computed efficiently. Once the eigenvalues of a matrix are known, calculating its exponential is extremely simple ($O(\sqrt{\text{Dim}(A)})$ in fact). Thus by taking advantage of the fact that the matrix is sparse allows the algorithm to terminate more quickly.

3 Project design

(More detailed documentation for the suite can be found in the code. Run `make documentation` in the repository's home directory and a directory containing technical documentation in HTML will be generated.)

This project has two overall goals: 1) to create code that will perform wave propagation as efficiently as possible, and 2) to make this implementation as portable and easy to use as possible to use as possible. To that end, there will be four functions implemented, each representing different approximation methods.

The primary computational structure used in the program is a struct `SqMat`. This struct is a square matrix of real numbers, and the current implementation of the the library only supports real matrices.

```
double* expv(struct SqMat mat, double t, double * v, double tau, double minerr);
```

A function to perform the the Krylov estimation for $\exp(t * mat)v$ where `tau` is the time step used and `minerr` is the minimum error allowed (must be greater than machine error on double operations). This algorithm is optimized for the case `mat` is large and sparse. (It is the user's responsibility ensure that $Dim(v) = \sqrt{Dim(mat)}$.)

```
struct SqMat expPade(struct SqMat mat, int p, int q);
```

An implementation of an older algorithm to implement the p,q-Pade approximation for $\exp(mat)$. By default, $p = q = 14$ should be used.

```
struct SqMat expSOFT(struct SqMat mat);
```

An implementation of the SOFT method of matrix exponentiation. This is not yet fully implemented.

```
struct SqMat expTaylor(struct SqMat mat, int k);
```

An implementation of the Taylor approximation of

$$\exp(mat) \approx \sum_{n=0}^k (mat)^n / n!$$

This should not be used in practice but is included for completion.

In addition there are a few other functions calculated in the the suite. The assist the above calculations, but are also documented and available to users as well. All of the following are standard in Matlab, but a pure C implementation outperforms Matlab by approximately factor 10 in runtime.

```
struct SqMat invSqMat(struct SqMat);
```

An implementation of Gaussian matrix inversion.

```
struct SqMat multSqMat(struct SqMat mat1, struct SqMat mat2);
```

An implementation of real matrix multiplication.

```
struct SqMat powSqMat(struct SqMat mat, int j);
```

A function to take the j-th power of matrix `mat`.

A small set of example programs demonstrating how the functions should be implemented in the near future.

4 TODO

The project now forks into two sub-project: 1) begin the GPU optimization process and 2) implement more features.

While GPU optimization is simple enough to understand, the practical steps are more finely detailed. As of now, we are tentatively considering using SPIR-V from the Vulkan API. The advantaged of this implementation over something like CUDA or OpenCL is that SPIR-V does not require the programmer to configure the the code for any possible computer architecture. SPIR-V allows the programmer to focus on higher level program design. One draw back on SPIR-V, especially in comparison to CUDA, is that the Vulkan API has not been generally accepted, so one older systems, the library may not run. (Vulkan is an API slotted to replace OpenGL, and SPIR-V is its GPU coding paradigm.)

The other direction, implementing more features, is more strait forward. First the structs `struct Complex` and `struct SqMatCplx` for complex numbers and complex matrices respectively. These will be used in a new implementation of the above functions. These new functions are absolutely necessary, as nearly all Hamiltonians are complex. (The decision to focus on real matrices first was to get a functioning implementation ready as soon as possible, even with limited functionality.) Of course once the functions are implemented, they will need to be GPU optimized as well, but this should be more strait forward using the real versions of the functions as a draft.

Finally, after both of these steps are completed, wrappers should be implemented so that these functions can be called in Python, Matlab, and Julia. In addition, additional scripts should be written to implement machine learning APIs. However the exact design of this final API is not set in stone.

Ongoing performance tests against the state-of-the-art are being performed.

5 Conclusion